



# **Advanced Document-Encoding and PDF-Generation Software to Improve Usability and Machine Readability of PDF Documents**

*Medical Imaging/Bioinformatics, Software Development,  
Academic Publishing, Manufacturing/Engineering, E-Commerce, and Real Estate*

## **INTRODUCTION**

We envision a much more expansive role/market for PDF files by providing advanced text encoding and document generation features. Such text-encoding and document-generation features would have a significant impact on the *interactive* capabilities of PDF files. As a result, the use of PDF files in everyday activities, from shopping to canvassing property listings to browsing cars for sale, would become much more commonplace. Equally so, the PDF format as used in professional/workplace activities, such as engineering and academic publishing, becomes more powerful as PDF files themselves become more interactive.

Currently, the relatively restricted adoption of PDF is driven, in part, by the fact that traditional PDF offers only limited options for interactivity, through features such as PDF scripting.<sup>1</sup> Another reason for the restricted use of PDF is that PDF generation tools are not programmed to take advantage of interactive features, thereby limiting what one can do with a PDF. Such limitations can, however, be rectified by using advanced text encoding that would enable one to embed metadata within PDF documents (metadata that identifies special actions to associate with individual PDF coordinate regions) and to implement callback handlers in PDF viewers, so as to support user actions and context menus linked to PDF coordinate regions.

PDF files augmented with these extra features can make the PDF experience much more interactive and user-friendly, allowing PDFs to play a larger role in application development. For example, PDF documents may be connected to external applications that provide graphics and multimedia or to application-specific features such as shopping carts and diagram/plotting engines. We envision such interactive PDF capabilities as having use cases in a broad number of vertical markets, such as e-commerce, real estate, manufacturing, bioinformatics, and publishing.<sup>2</sup> Here are three use-cases below:

## **REAL ESTATE USE CASE: PROVIDING VIRTUAL TOURS WITHIN A PDF VIEWER ITSELF**

As a concrete example, real estate listings could be provided in a PDF format which would embed virtual reality/panoramic photography viewers built via technologies such as Matterport. Our tools would then provide a GUI component that would be included in a PDF viewer, providing access to the panoramic photography that would make the tour a much richer experience. For instance, the PDF engine could allow realtors to furnish the information which users see when they click on Matterport tags that are attached to specific locations on the property. This would allow realtors to convey details such as sun exposure in each room, closet size, and cabinet arrangement and other features of the kitchen. In a browser-based viewer, in contrast, the realtor

---

<sup>1</sup> In fact, scripting itself is limited because many PDF users/viewers disable scripting for security reasons.

<sup>2</sup> More generally, interactive PDF technology can integrate PDF viewers with a wide range of software applications. This interoperability can be achieved by encoding and sharing the textual content from which a PDF document is generated, so that this raw text is available to user-action handlers within the PDF application, and by annotating external data and code which may be linked to the PDF application so as to cross-reference these external resources with the document text.

would have no means of customizing what information the user would find via clicking on tags, because in a web browser, unlike within a specialized PDF viewer, Matterport tags can only be connected to external websites, such as those of appliance manufacturers. In short, by providing a virtual tour *within* the PDF viewer itself, rather than on an external website/browser-based viewer, the PDF application may detect (via application signals) which tags the users are selecting and respond with spot-on information, which could make the difference with keeping the user interested in the property and eventually going forward with a purchase offer. At the very minimum, having all this information associated with locations in the virtual tour would keep the user engaged for a longer period of time, which may inspire the user to visit the property in person.

## E-COMMERCE USE CASE: MAKING PURCHASES THROUGH THE PDF APPLICATION ITSELF

The way this would work is that digital product catalogs would have an embedded shopping cart so that the purchase could be made *through* the PDF application itself instead of via a web portal. This is important because it would allow users, upon checking out, to have the full benefit of the product specs and images (including 3D graphics and multimedia) at that time rather than having to scroll back to get the product specs. This would keep users focused on the catalog item(s) and minimize the likelihood of users abandoning their shopping cart to scroll back to the product specs while failing to return to the shopping cart to check out.

## TECHNOLOGY STACK: ACADEMIC PUBLISHING AS A USE CASE

The important programming layers for generating interactive PDF files, and the relationship between text encoding and data/code annotation, are outlined below. This same technology applies to multiple vertical markets, but we are pointing to academic publishing in our description of the technology:

1. **A New Text Encoding Protocol** — Introducing a novel text encoding for document markup, including PDF generation, which produces user-friendly and machine-readable text representation and metadata that may be embedded as a PDF attachment. The resulting supplemental material, which a compatible PDF viewer can automatically extract, would consist of two parts: (1) a direct encoding of document text in the form of a character stream with a fixed character size (e.g. one byte per character); and (2) metadata representing the character positions/ranges and PDF coordinates of document content such as sentences, paragraphs, citations, keywords, and figures/graphics.

With respect to document text, having a character stream (or several character stream “layers”) as a direct representation of textual content — rather than “scraping” characters from the PDF directly — leads to more accurate searching and UI features (e.g. copy/select). For example, with one single context menu action users can copy an entire sentence/paragraph with just one click. This is made possible because the character streams employ a customizable encoding (instead of a generic format like Unicode) so that the document can specify ahead of time how text segments should be translated into character sequences for searches or selection-copy features; notably, how to handle foreign letters, equations, footnotes, and so forth, details which ordinarily complicate basic functionality such as “copy a range of text to the clipboard.” Fixed-width character encoding ensures that all text segments can be identified by simple numeric ranges and that text can be decoded by iterating over one character at a time in isolation (in contrast to Unicode multi-bytes or XML character entities). With respect to metadata, supplemental files encode annotations which map character-stream positions to PDF coordinates for document content that would be a target for user interactions, such as sentence boundaries (to enable automatic sentence-copy features), figures/graphics (for context menus providing figure information), citations (for capabilities such as searching cited publications’ DOI) and so forth. We have put together a demo PDF viewer which reads the character and metadata files



and uses this data to provide new capabilities that are not found in other PDF programs (including Acrobat).

2. **Dataset Integration and “Microcitations”** — When integrating data sets with PDF documents, our text-encoding protocol is able to support fine-grained cross-references between individual parts of the data set and individual locations or regions of the document. This contrasts with the conventional approach wherein connections between data sets and publications are noted only coarsely, such as when a “Supplemental Materials” section at the end of a paper provides a link to download an accompanying data set, but only with a cursory overview of what is inside the data set. By contrast, my team envisions granular cross-references where, for example, a concept embodied in a data set (for instance, a statistical parameter instantiated by a table column or datatype field) can be mapped to the sentence in the publication where that concept is defined or is first introduced. Conversely, diagrams or figures could be linked from a document to the relevant part of an accompanying data set, such as individual tables which generate a statistical plot or to source code files where statistical distributions are calculated. Thus, in order to support these granular cross-references — which are a form of “microcitation” as this term is used in science — we have prototyped novel “dataset creator” tools that build customized “dataset applications” for viewing and analyzing dataset contents.

In the spirit of initiatives such as Research Objects and FAIRsharing (Findable, Accessible, Interoperable, Reusable) — standards which have been prioritized by institutions such as the NIH, the Bill and Melinda Gates Foundation, and the Chan Zuckerberg Initiative — dataset applications using our LTS tools would combine code, data, and visualization/GUI components and, moreover, would be launched directly from the PDF viewer itself. A standardized protocol cross-referencing PDF files with data sets (annotating specific parts of the document and the data set which are conceptually related) could also include specifications for launching external applications to view embedded multi-media content, such as 3D graphics or video files. In effect, support for dataset annotations could be merged with an annotation system whereby files embedded in a PDF document — potentially linked to visible PDF elements such as a 2D figure representing a view onto a 3D model or a still frame from a video — are annotated with metadata allowing the PDF viewer to export resources to an external application. Integrating PDF viewers with multi-media content and datasets can be important because, at present, there is no software platform where datasets and their corresponding publications can be seamlessly integrated. One is, therefore, less likely to find (and no less cite) the article associated with the dataset. In contrast, by using tools that effectively integrate datasets with their corresponding publications this would increase the chance of both finding and citing the publication. In addition, by using these document preparation tools, as described above, data would be much more meticulously curated resulting in a more rigorous presentation of the data itself, undoubtedly making it more appealing to researchers to cite such data, along with the journal article or book wherein the data is contained, when presenting their own work.

3. **Cloud Hosting** — At present, data-hosting sites (such as Dryad, Dataverse, or Open Science) and document repositories (such as Science Direct, Springer Nature, or ScienceOpen) usually operate in parallel, but a next generation of content providers (combining publishing and data-hosting platforms) can host both data and documents simultaneously, in an integrated fashion, allowing for searches across both publications and data sets. This means that data sets have to be indexed and searchable as well. When a document is linked to a data set, the document viewer (canonically a PDF viewer) would then have the option of accessing the document source and retrieving dataset information — for example, identifying software prerequisites for using the data set and explaining to readers how they can acquire and access the dataset contents — and once the data set is downloaded and the relevant software is identified (either as an application embedded in the data set or as an external program) the PDF viewer could then connect with that software component as part of its user-interaction interface. These kinds of workflows would require APIs and multi-application networking protocols that could be concretized by providing cloud-hosting services or projects such as a reusable docker container, which demonstrate API and protocol implementations as well as document and dataset hosting capabilities.

